

Unpicking the weave:

**A practitioners view of challenges faced in
debugging multithreaded/ many-core applications**

**Neil Hickey
Samsung Research
Institute: UK**

A decorative graphic at the bottom of the slide features several overlapping, translucent blue wavy lines that flow from left to right, creating a sense of movement and depth.

Samsung Research UK

- ❖ **Research Institute focused on developing innovative software solutions for mobile, web and smart TV.**
- ❖ **SRUK are heavily involved in Khronos standards group, leading the UK Khronos chapter and pushing development of open standards for graphics/ compute APIs**
- ❖ **SRUK specializes in optimisation and debugging of Android software utilising the CPU and GPU (OpenGL ES, OpenCL)**

Multithreaded debugging

❖ Why multi-threaded.

- **SRUK optimise applications making use of the GPU. This is typically a single application utilising the GPU and multiple threads of the CPU simultaneously.**

❖ Multicore / Many core

- **Debugging multi-threaded CPU code is difficult enough.**
- **Debugging GPU code is an order of magnitude more difficult**

Multi-threaded development challenges

❖ It is difficult.



Problems of parallel programming

❖ **Data races**


❖ **Deadlock / Live lock**

❖ **Serialisation**

❖ **Simultaneous memory updates**


Data races

- ❖ Multiple threads attempting to access the same piece of data.
- ❖ Non parallel behaviour while another thread is updating.




Deadlocks

- ❖ Multiple threads attempting to secure a lock.
- ❖ Two threads each capture a lock on a resource and freeze up attempting to capture the lock on the other threads resource



Live locks

- ❖ Two locks needed, one thread gets one, another thread gets another. Each thread tries to get the other and when it fails, frees its own lock.



Serialisation

- ❖ Multiple threads accessing shared resource, serialising on that shared resource therefore running no better than synchronous.



Simultaneous memory updates

- ❖ Multiple threads accessing data, e.g. Free pointers and causing data corruption.



Multi-threaded debugging tips and tricks

- ❖ Trace buffers
- ❖ Debugging optimised code
- ❖ Use asserts and capture core files
- ❖ Perform object hand-over rather than shared objects
- ❖ Built-in debug support
- ❖ RAI for locking/unlocking mutexes
- ❖ Conditional breakpoints

Trace buffers

- ❖ Instrument interesting sections of code with trace buffer writes to allow a textual representation of the access order.

Example

```
int Thread1()
{
    ...
    TraceEvent(beforeMsg);
    doSomething();
    TraceEvent(afterMsg);
    ...
}

int Thread2()
{
    ...
    TraceEvent(beforeMsg);
    doSomethingElse();
    TraceEvent(afterMsg);
    ...
}
```

Trace buffers example output

```
Trace 1 : StartWork Thread 1   time 79ms
Trace 2 : EndWork Thread 1     time 83ms
Trace 3 : StartWork Thread 2   time 94ms
Trace 4 : EndWork Thread 2     time 104ms
Trace 5 : StartWork Thread 1   time 111ms
Trace 6 : StartWork Thread 2   time 112ms
Trace 7 : EndWork Thread 2     time 125ms
Trace 8 : EndWork Thread 1     time 127ms
```

Potential problem here where Thread 2 has overtaken Thread 1, which may be unexpected

Debugging optimised code

- ❖ Disable instruction reordering
- ❖ Drop optimisation level
- ❖ Implement logging rather than single stepping code.

Use asserts and capture core files


- ❖ Detect memory corruption by using asserts
- ❖ Captured core file can be run in the debugger to examine state

Perform object hand-over rather than shared objects

- ❖ Manage object ownership explicitly rather than allowing all threads to access shared memory.
- ❖ Single thread owner and destroyer.

Built-in debug support

- ❖ Produce your own application specific logging to assess state of failing applications after the fact



RAII for mutex locking unlocking

- ❖ For functions that can throw assertions or have many exit points, utilise RAII for mutex locking and unlocking.

C++11 Example

```
int Thread1()  
{  
    MyLockingClass<std::mutex> myLock(gMutex);  
    ...  
} // MyLockingClass destructor is called here irrespective of  
   how the function exited, releasing the mutex
```

C Example using GCC extension

```
int Thread()  
{  
    RAII_VARIABLE(MyMutex, mutex, lock(mutex), unlock());  
    ...  
}
```

Conditional breakpoints/ ignore

- ❖ Only stop the process if a preset condition evaluates to true at the time the breakpoint is hit.
- ❖ Ignore the next X number of breakpoint fires before stopping

Multi-threaded debugging shortcomings

- ❖ Single stepping multiple threads
- ❖ Viewing thread state
- ❖ Viewing data
- ❖ Reproducibility
- ❖ Architectural knowledge

Single stepping multiple threads

- ❖ Stepping on one thread continues all other threads.
- ❖ Getting around this involves setting breakpoints everywhere
- ❖ This is difficult to manage and not easy to set up in the first place

Viewing thread state

- ❖ How to represent thread state simultaneously.
- ❖ Requires to manually change active thread in gdb.
- ❖ Error prone if you step on the wrong thread.

Viewing data

- ❖ How to view data on different threads simultaneously.
- ❖ Requires thread to be changed explicitly
- ❖ Difficult to visualise data for very large numbers of threads in this way

Reproducibility

- ❖ How to reliably reproduce a failing state
- ❖ Capturing thread affinity
- ❖ Reapplying thread affinity

Architectural knowledge

- ❖ Heap allocation

- ❖ Stack allocation

- ❖ Shared object data access between processes

Things that need to be improved

- ❖ Single step/next will stop when any thread hits the next instruction / source line.
- ❖ Getting and setting the core affinity for threads spawned.
- ❖ Viewing thread local data across all threads.
- ❖ Automatic trace buffer insertion.

Thank you.

A decorative graphic at the bottom of the slide consisting of several overlapping, translucent blue wavy lines that flow from left to right, creating a sense of movement and depth.