



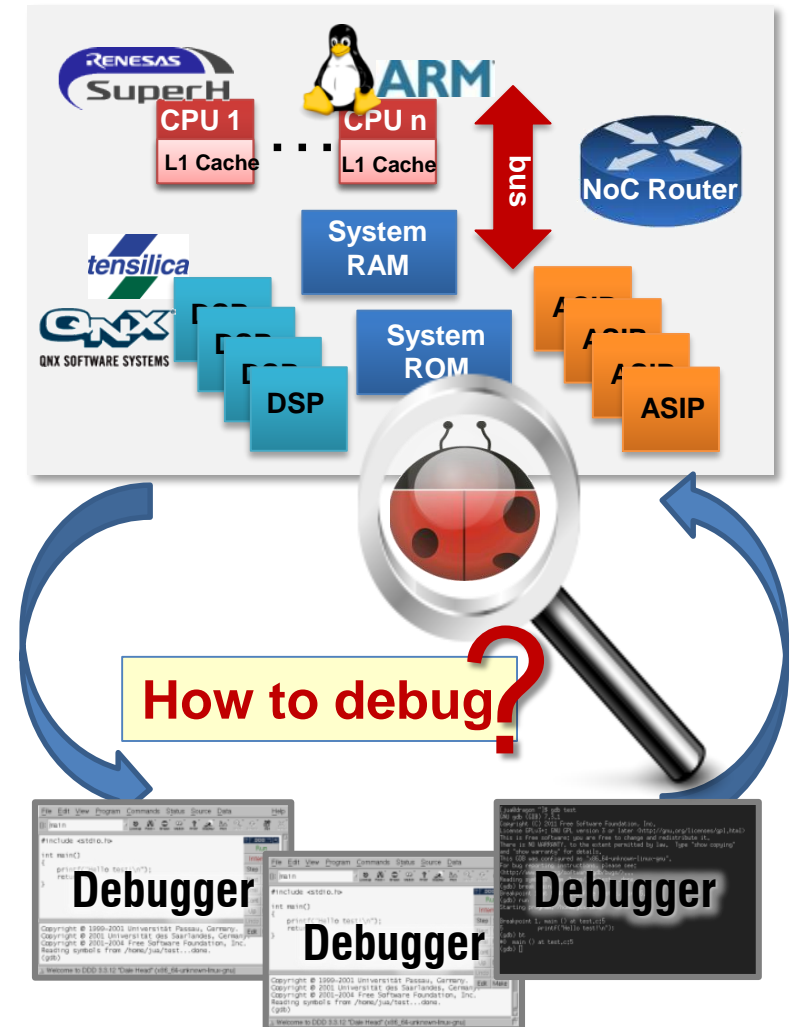
Automatic Exploration of SW Concurrency Bugs through Deterministic Behavior Control

Luis Gabriel Murillo, Rainer Leupers

MAD Workshop
14.11.13, Munich, Germany



- **MPSoCs**
 - Complex communication
 - Shared memory, *KPN* and *SDF* models, message passing...
 - Co-existing OSs, middle-wares...
- **Concurrency → Non-determinism**
- **Many-cores → Many debuggers?**



- **MPSoCs are non-deterministic**

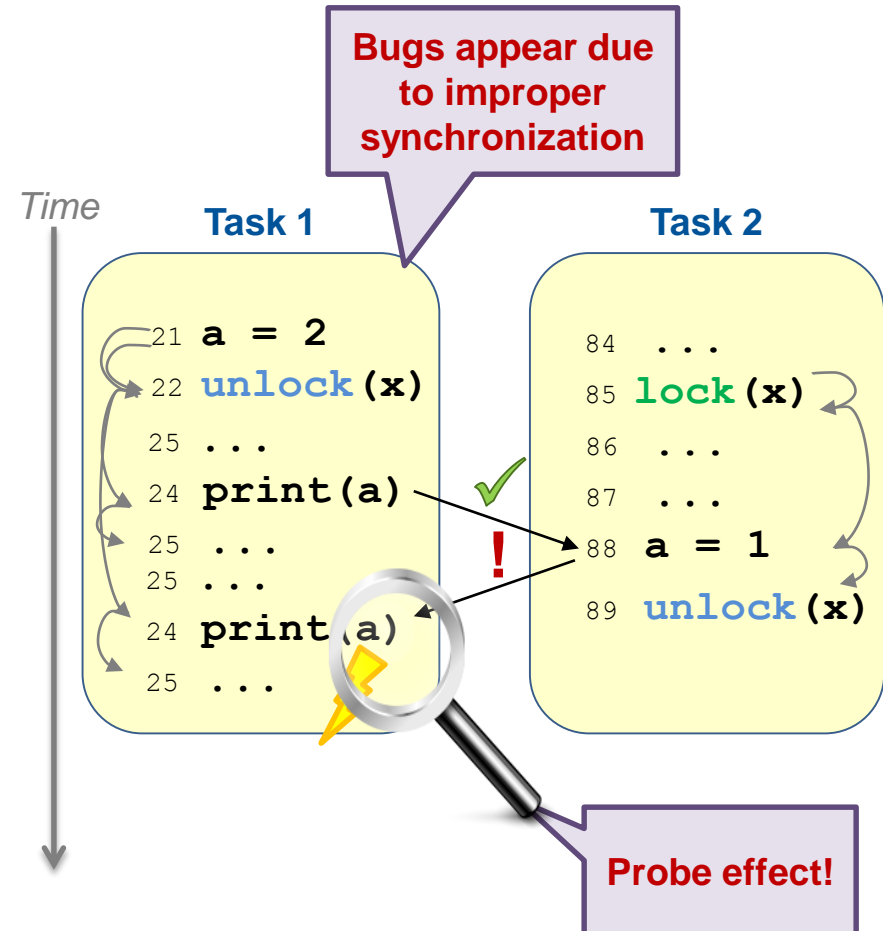
- **Concurrency Bugs**

- Races (order and atomicity violations)
- Deadlocks, livelocks...

- **Difficult to:**

- Find
- Understand
- Reproduce

→ Remain **unnoticed**



MPSoC Debug Challenges

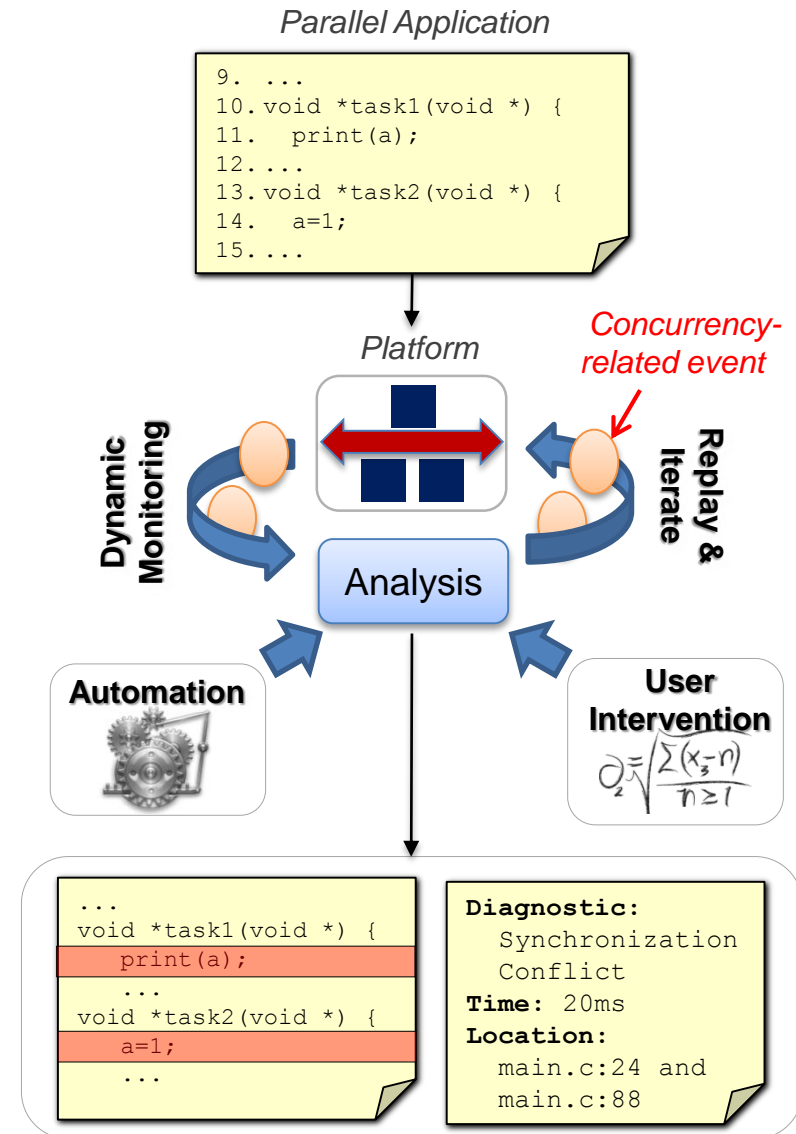
➔ Methodology Overview

Event-based Debugging

Determinism Analysis & Behavior Control

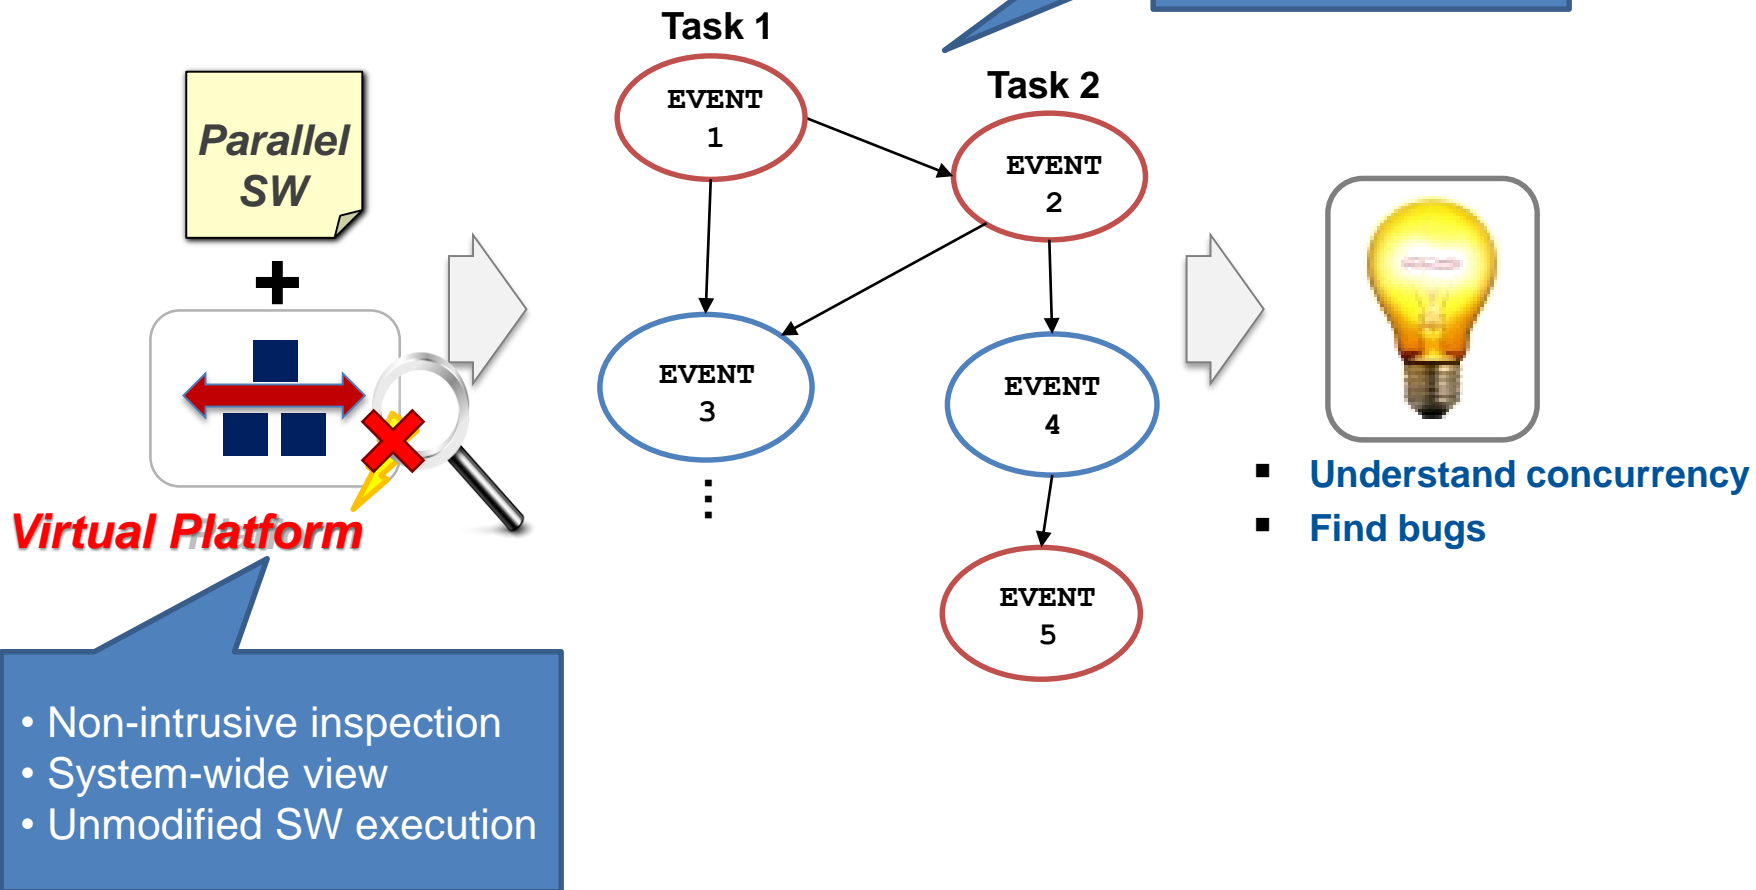
Results and Conclusions

- **Goals:**
 - Help in finding concurrency bugs
 - Unique methodology / debugger for different platforms
 - Tool for SW programmer
- **Key aspects:**
 - Abstraction
 - Automation
 - Retargetability
 - Scalability



- **Abstracting away program flow:**
 - Focus on programmer level actions / concurrency related events

All synchronization, task management, message passing, shared memory...



	AVIO (Lu et al. '06)	Chess (Microsoft '08)	Portend (EPFL '12)	This work
Target system	x86	Windows	LLVM	Virtual Platform
Target application	C(++)	.NET	Pthread	SW + HW
Non-intrusive	X Instrumentation	X Wrapper	X Symbolic execution	✓
Deterministic replay	X	✓	✓	✓
Deterministic program exploration	X	✓	✓	✓
Extensibility	X	X	X	✓

MPSoC Debug Challenges

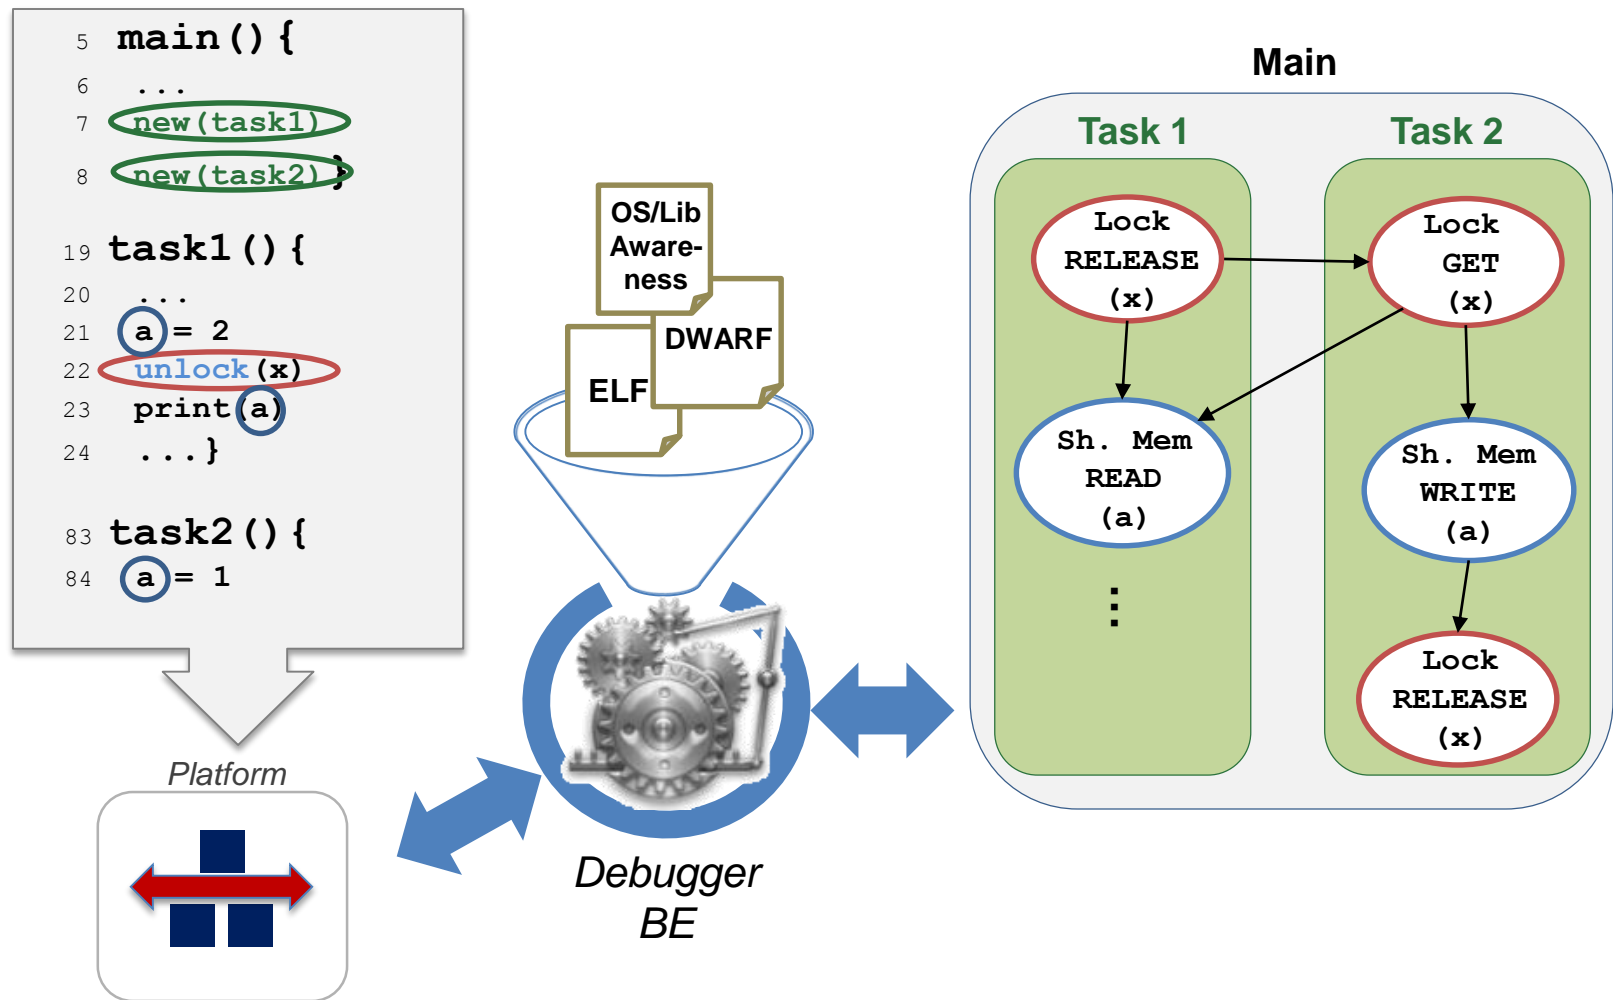
Methodology Overview

→ Event-based Debugging

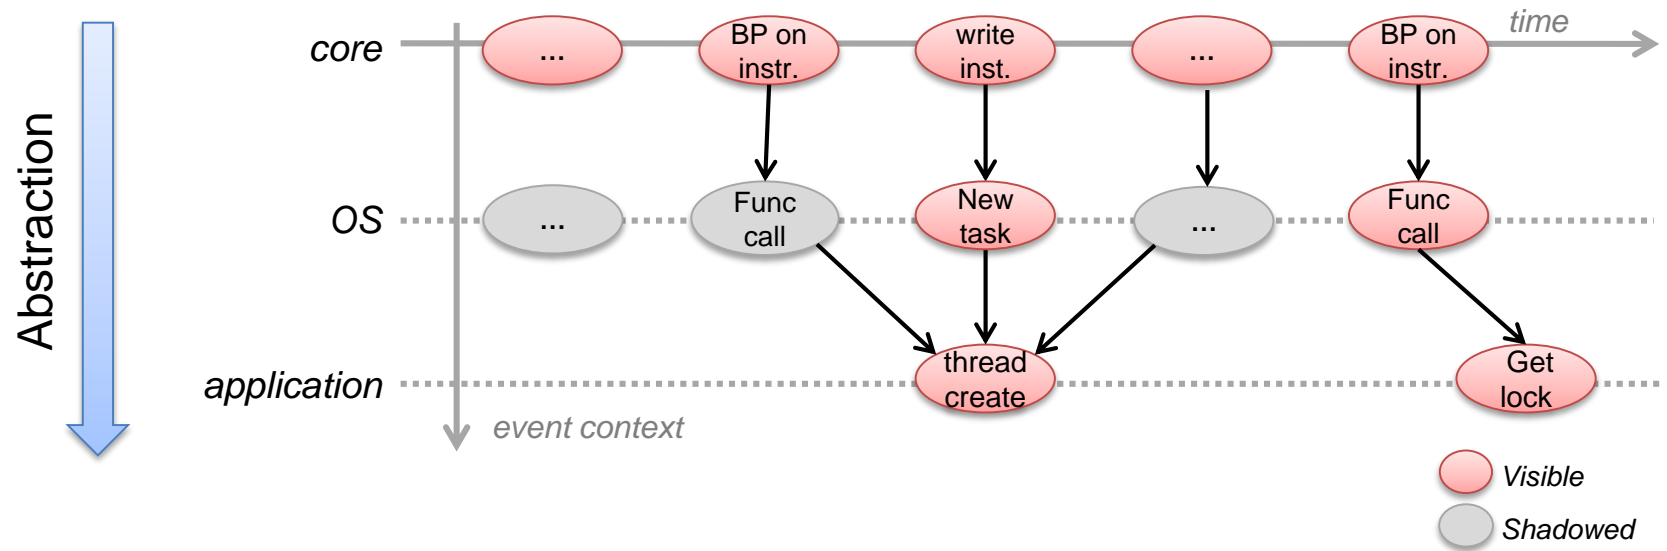
Determinism Analysis & Behavior Control

Results and Conclusions

■ Debugger framework for **Dynamic Monitoring**



- **Problem:** High-level atomic events for analysis but fully trackable to origins
- **Solution:**
 - Bi-dimensional **composition**: *time, context*
 - Propagation of **semantic** information



- **Reveals the order** of programming-level events
 - “Understanding” the application
- **Identification of relevant source code location / task / core**
 - Dynamic monitoring with source debugger
- **No source code instrumentation, no changes to target SW, non-intrusive monitoring...**
- **Trace captures one single execution**
 - One single “task interleaving” 😊
 - Other possible interleavings? ☹️

MPSoC Debug Challenges


Event-based Debugging

Bug-pattern Assertions

➔ Determinism Analysis & Behavior Control

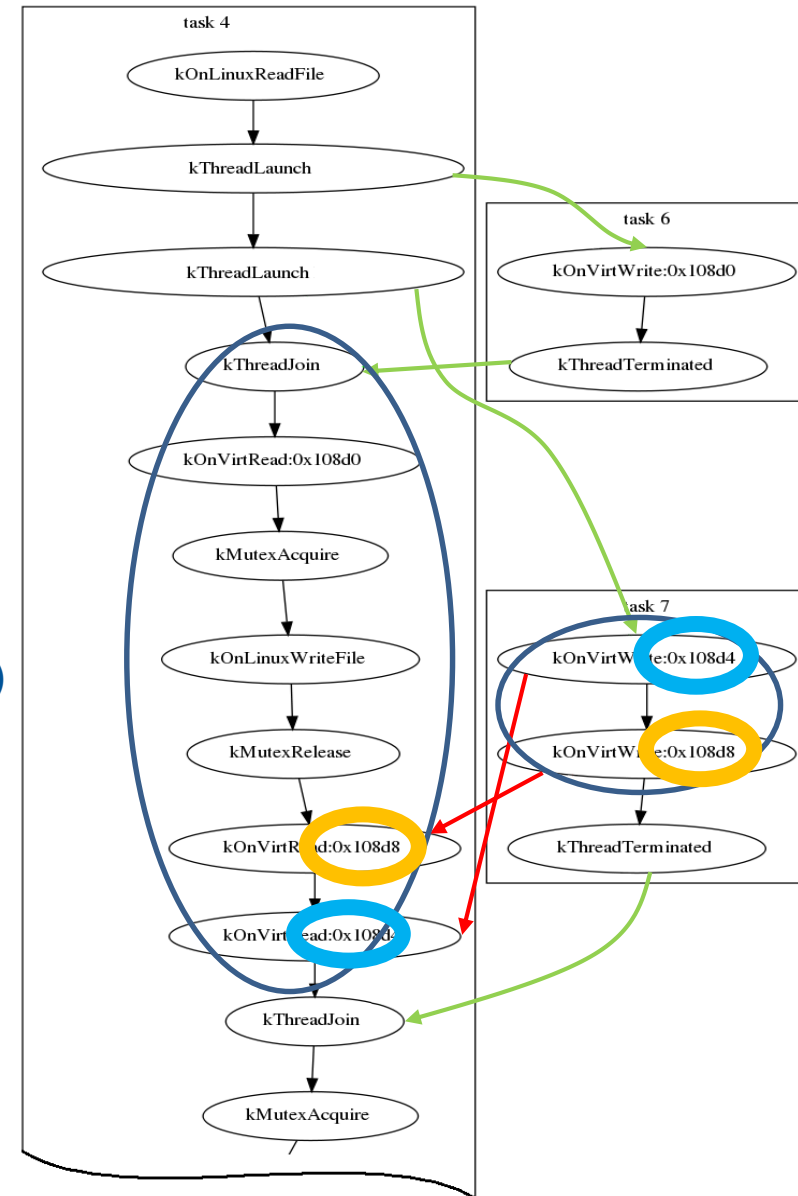
Results and Conclusions

- **Problem:** “One single execution is not enough to spot concurrency bugs”
- **Solution:** **concurrency analysis** and **controlled replay**
 - Investigate suspicious interleavings
 - Identification of **non-determinism** ‘with notable effect’
 - Provoke bugs which are hidden! 😊

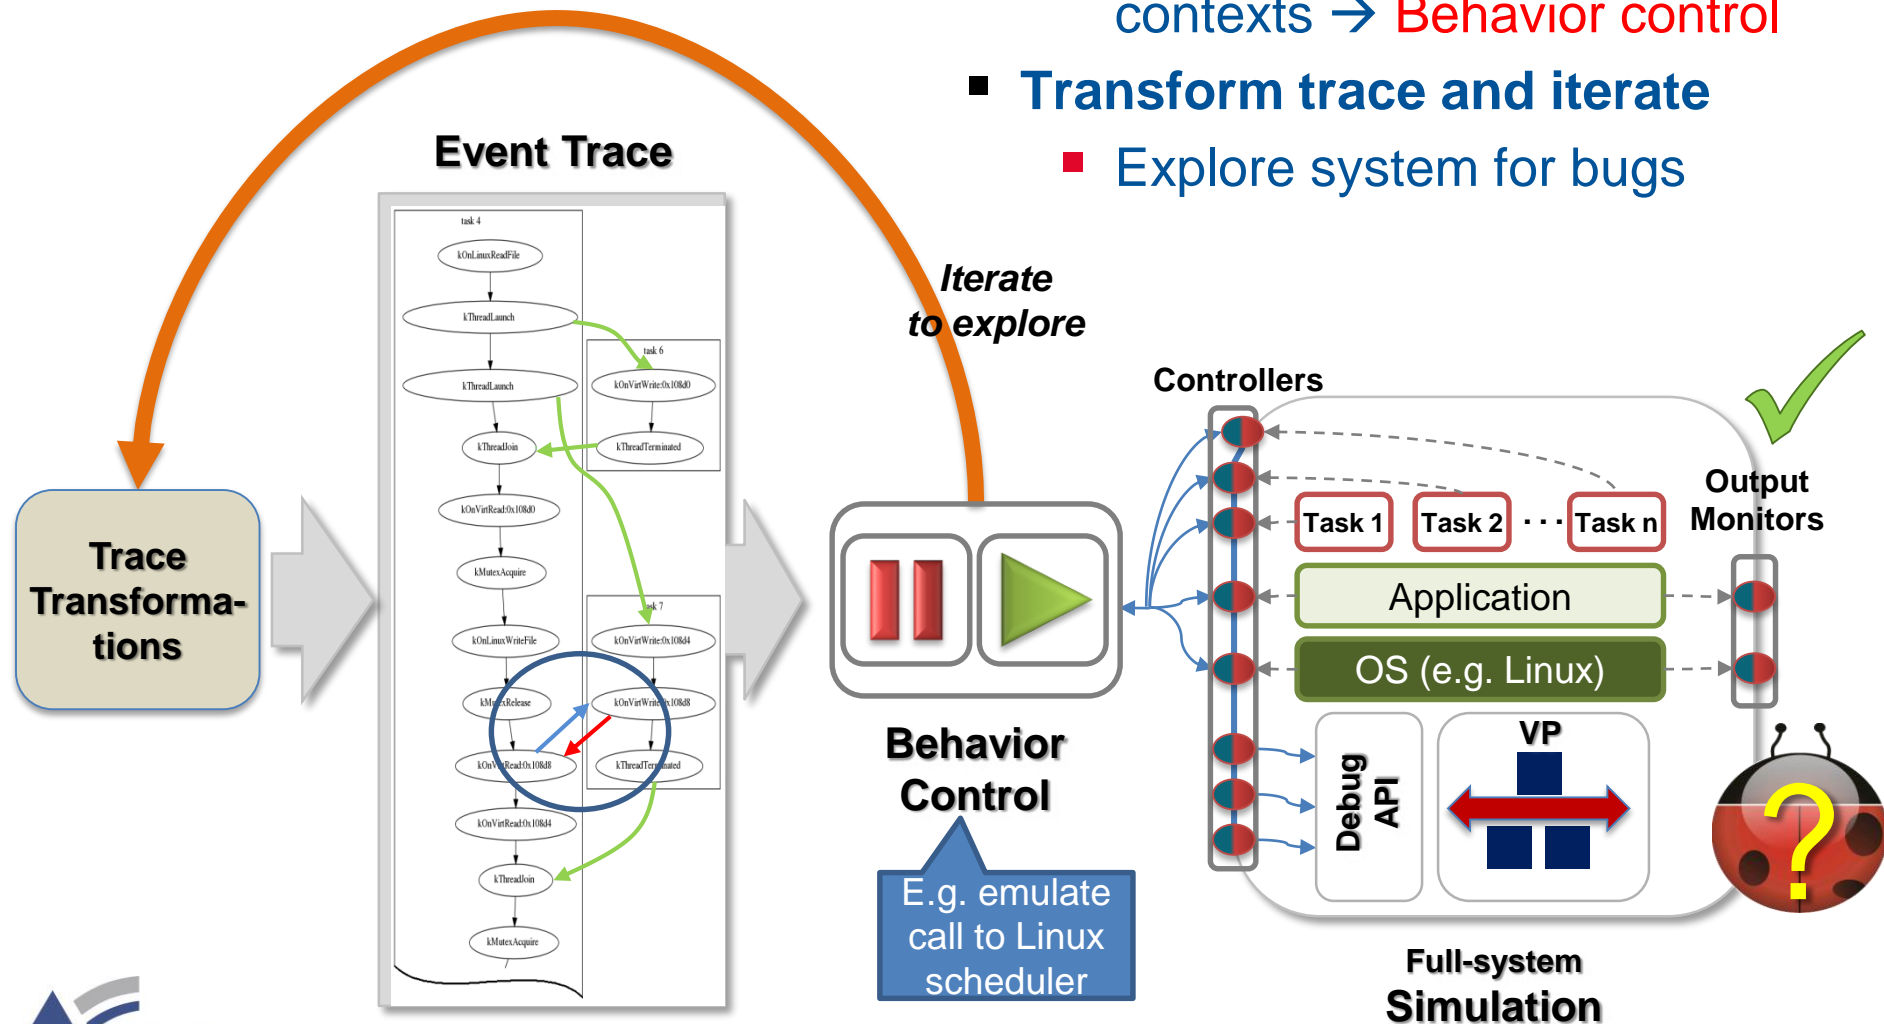


■ Concurrency analysis and conflict extraction:

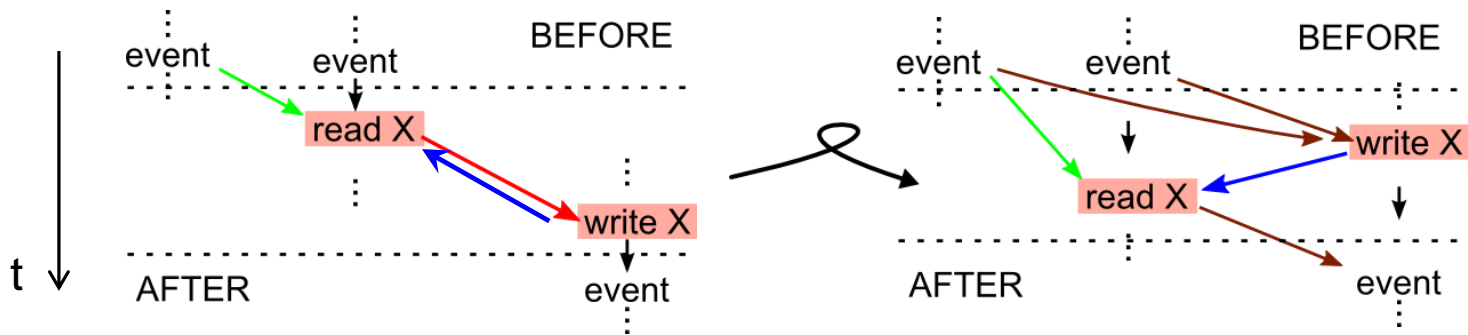
1. Identify synchronization
 - Mark “**always happen**” event orders (“happens before” analysis)
2. Identify “**always concurrent**” events
3. Identify event **dependencies**
 - On shared resources (“Visit/Modify”)
4. Identify **conflicts**
 - Dependencies not in sync
5. For exact replay or bug provoke:
 - Enforce order of conflicting events
 - Minimal set of event pairs



- **Event-based replay**
 - Suspend/resume event contexts → **Behavior control**
- **Transform trace and iterate**
 - Explore system for bugs



- **Swapping a conflicting event order**
 - Locally invert a constraint
 - Single swap is **safe** and likely to **change behaviour**
- **Swapping a constraint**
 1. **Swap** event pair order
 2. Add **repair constraints** for locality



→ Random Constraint Swapping

MPSoC Debug Challenges

Event-based Debugging

Bug-pattern Assertions

Determinism Analysis



Results and Conclusions



- **EURETILE** (www.euretile.eu)

- European reference tiled architecture experiment
- Many-tiled system for embedded and HPC



- **Multi-core Synopsys Virtual Platforms**

- **ARM Versatile Express** with 4 Cortex A9
 - SMP Linux 3.4.7, pthreads, SPLASH-2

Results ARM Versatile Express

Event-based Framework

	Retargetable BE	High-level Monitors
Adaptation Effort	~1 man-month	~2 man-days

Monitoring and Analysis

	Synthetic	SPLASH-2
Total events (no SM)	~500	600 – 123k
Total events	~2500	3000 – 1.9M
Overhead	~3x	~3x (WC:60x)
Replay Constraints	~50	500 - 3200

→ Event trace and analysis results

	Filtered conflicts			
	Total	Sync ■	Mutex ■	Conflict ■
Count	284	260	23	1
rel.		91.5 %	8.1 %	0.4 %

→ Unsynchronized dependency in OCEAN event trace


- Variable at 0x72014: `global->psibi`

```

516: /*LOCK(locks->psibilock)*/
517:  global->psibi = global->psibi + psibipriv;
218: /*UNLOCK(locks->psibilock)*/
    
```

```

item0: previous modify (6) at 1405
(6, kNone).kOnVirtWrite(0) @00072014
@000199dc: slave1.C:517
===
item1: current visit (4) at 19913
(4, kNone).kOnVirtRead(0) @00072014
@000199dc: slave1.C:517
    
```



E.g., Result of Exploring Bugs in OCEAN

Ocean simulation with W-cycle multigrid solver

Processors : 2
Grid size : 6 x 6
Grid resolution (meters) : 20000.00
Time between relaxations (seconds) : 28800
Error tolerance : 0.1



iter 4, level 0, residual norm 9.69823850e-02, work = 4.000
iter 1, level 0, residual norm 4.66194437e-04, work = 1.000
iter 1, level 0, residual norm 9.32388873e-05, work = 1.000
iter 1, level 0, residual norm 2.76407790e-04, work = 1.000
iter 1, level 0, residual norm 5.52815581e-05, work = 1.000
iter 1, level 0, residual norm 1.07756867e-03, work = 1.000
iter 1, level 0, residual norm 2.15513732e-04, work = 1.000
iter 1, level 0, residual norm 1.04034932e-03, work = 1.000
iter 1, level 0, residual norm 2.08068122e-04, work = 1.000
iter 1, level 0, residual norm 7.44197648e-03, work = 1.000
iter 1, level 0, residual norm 1.48839561e-03, work = 1.000
iter 1, level 0, residual norm 7.02009090e-03, work = 1.000
iter 1, level 0, residual norm 1.40399094e-03, work = 1.000

Ocean simulation with W-cycle multigrid solver

Processors : 2
Grid size : 6 x 6
Grid resolution (meters) : 20000.00
Time between relaxations (seconds) : 28800
Error tolerance : 0.1

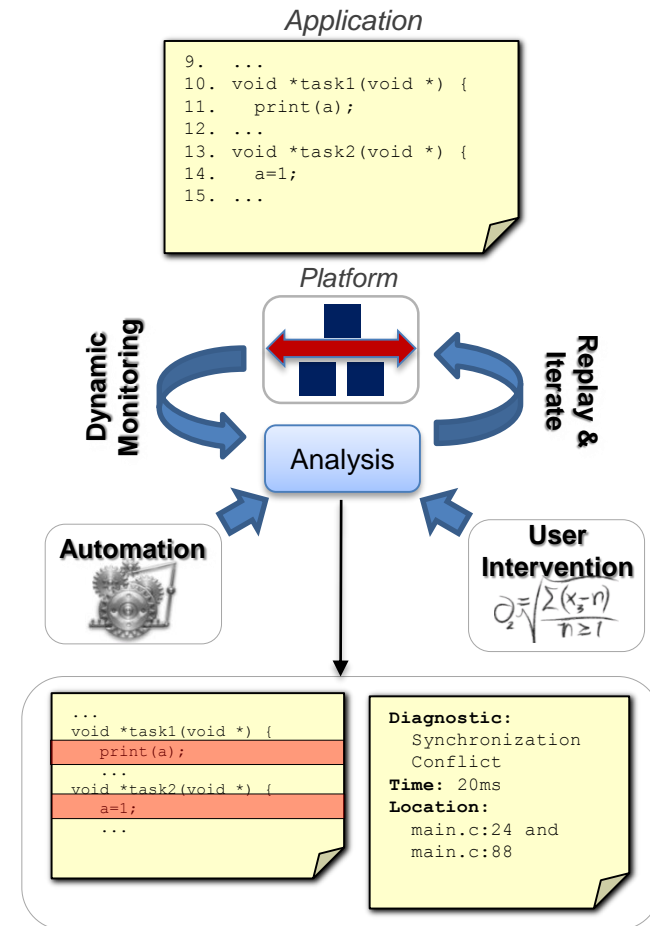


iter 4, level 0, residual norm 9.69823850e-02, work = 4.000
iter 1, level 0, residual norm 4.66194437e-04, work = 1.000
iter 1, level 0, residual norm 9.32388873e-05, work = 1.000
iter 1, level 0, residual norm 2.76407790e-04, work = 1.000
iter 1, level 0, residual norm 5.52815581e-05, work = 1.000
iter 1, level 0, residual norm 1.07756867e-03, work = 1.000
iter 1, level 0, residual norm 2.15513731e-04, work = 1.000
iter 1, level 0, residual norm 1.04035803e-03, work = 1.000
iter 1, level 0, residual norm 2.08068122e-04, work = 1.000
iter 1, level 0, residual norm 7.44197490e-03, work = 1.000
iter 1, level 0, residual norm 1.48839561e-03, work = 1.000
iter 1, level 0, residual norm 7.02022805e-03, work = 1.000
iter 1, level 0, residual norm 1.40399090e-03, work = 1.000

```
src/RandomSwapBugFinder.cc:299 : bug occurs when events happen in this order:  
first event: 0xc170f508 (4,kNone).kOnVirtRead(0) @00072014  
@000199bc: slave1.C:517  
second event: 0xc1702d48 (6,kNone).kOnVirtWrite(0) @00072014  
@000199dc: slave1.C:517
```

→The bug was found after one iteration.

- **MPSoC debuggers should:**
 - Facilitate intuitive ways to catch and identify system-wide bugs
 - Explore different concurrent interleavings
- **VPs + Concurrency Analysis → Good recipe to deal with concurrency bugs**
- **ICE's event-based debugging:**
 - Retargetability
 - Abstraction
 - Automation
 - Scalability





**Thanks!
&
Questions?**

