


# Observation of Multi-Core SoCs

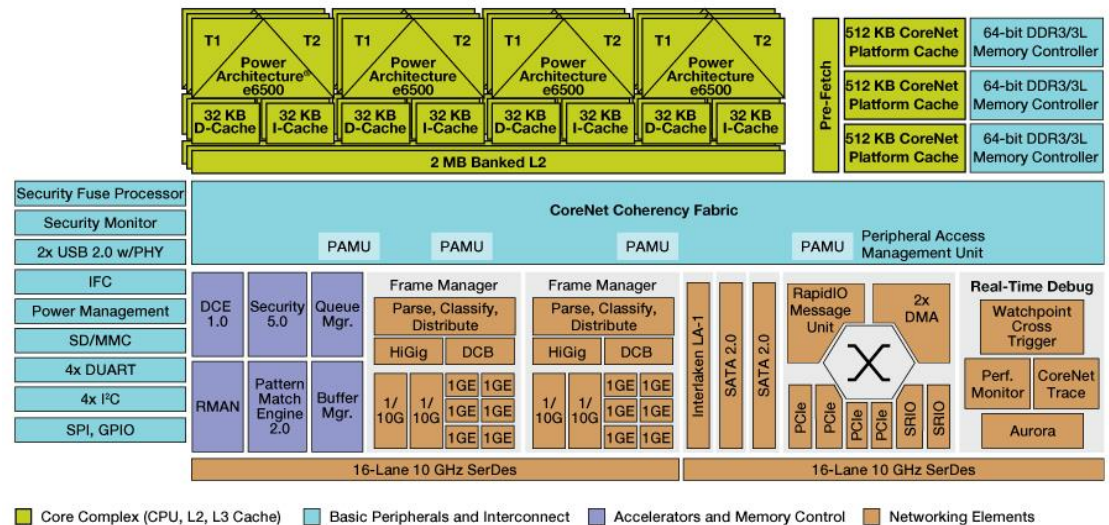
Alexander Weiss

Garching, 15.11.2013

Accemic GmbH & Co. KG



QorIQ T4240 Communications Processor

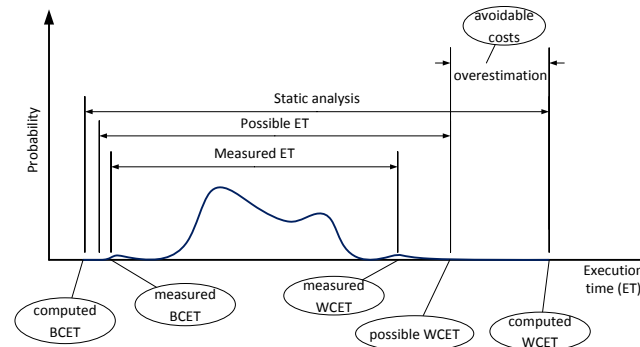

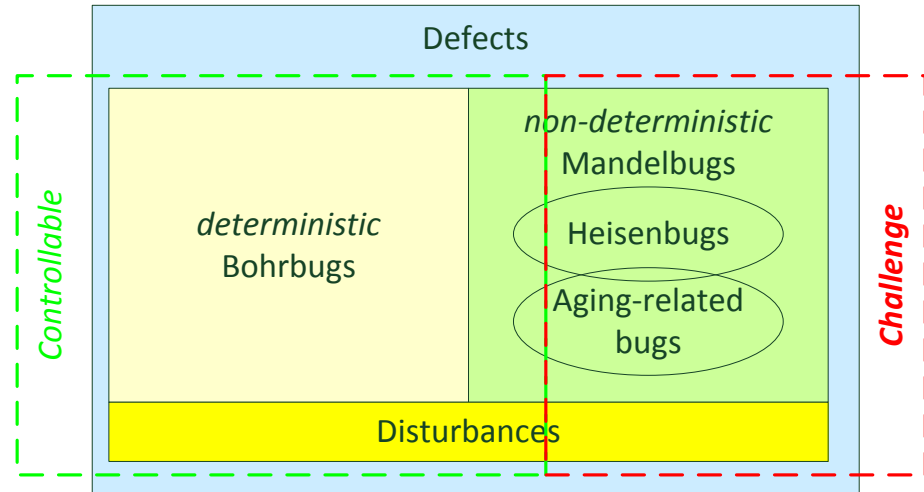


## You can design it, but can you debug it?

[Martin, Grant; Mayer, Albrecht; "The challenges of heterogeneous multicore debug ", Design, Automation & Test in Europe Conference & Exhibition (DATE), 2010]

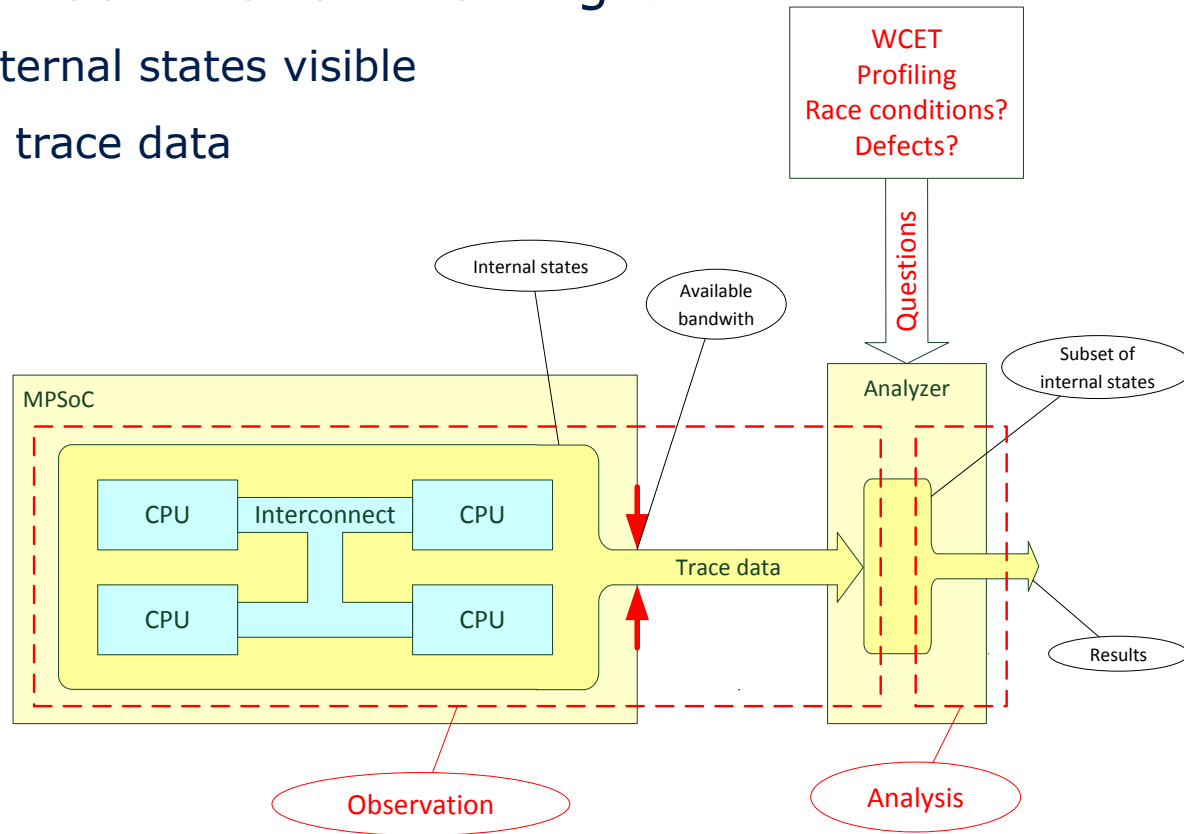
## Observation for

- Debug
- Tests / coverage analysis
- WCET measurement
- Detection of race conditions
- Profiling / optimization



## Multi-Core observation challenges

- Make internal states visible
- Analyze trace data



## What we want to know

- CPUs
  - Executed instructions
  - Clock cycles / instruction
  - Data access (value, address, direction)
  - Cache
  - CPU register
  - Events
- Bus system / Bus master peripherals
  - Bus master data access (value, address, direction)
  - Events (timeouts, concurrent access, splitted transfers, errors)
- Vector clock
  - Temporal assignment of CPUs and bus master operations

Shared resource	Mechanism
System bus	Contention by multiple cores Contention by other device - IO, DMA, etc. Contention by coherency mechanism traffic
Bridges	Contention by other connected busses
Memory bus and controller	Concurrent access
Memory (DRAM)	Interleaved access by multiple cores causes address set-up delay Delay by memory refresh
Shared cache	Cache line eviction Contention due to concurrent access Coherency: Read delayed due to invalidated entry Coherency: Delay due to contention by coherency mechanism read requested by lower level cache Coherency: Contention by coherency mechanism on this level
Local cache	Coherency: Read delayed due to invalidated entry Coherency: Contention by coherency mechanism read
TLBs	Coherency overhead
Addressable devices	Overhead of locking mechanism accessing the memory I/O Device state altered by other thread/application Interrupt routing overhead Contention on the addressable device - e.g. DMA, Interrupt controller, etc. Synchronous access of other bus by the addressable device (e.g. DMA)
Pipeline stages	Contention by parallel hyperthreads
Logical units	Contention by parallel applications
	Other platform-specific effects, e.g. BIOS Handlers, Automated task migration, Cache stashing, etc.

**UNDESIRED MECHANISMS AFFECTING THE TEMPORAL DETERMINISM**  
 from: Kotaba et al, "Multicore In Real-Time Systems – Temporal Isolation Challenges Due To Shared Resources", 2013

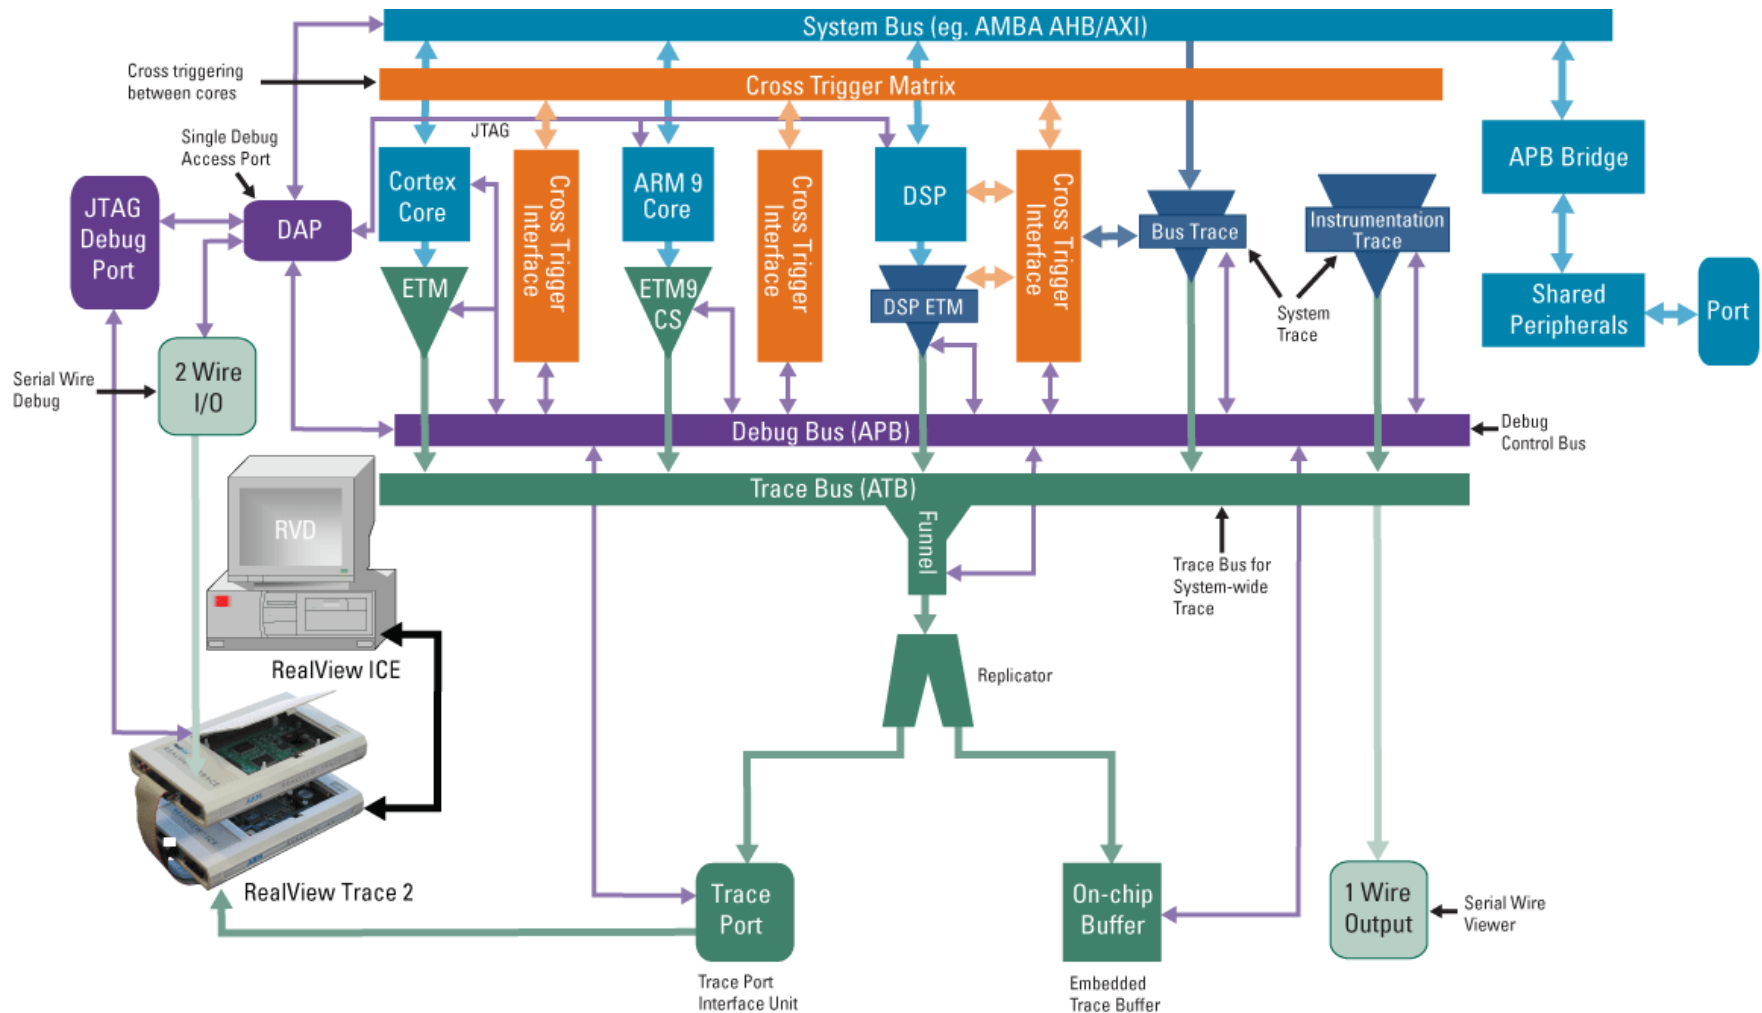
## Other requirements

- Real-time capability
- Non-intrusiveness
- Concurrent observation of multiple CPUs / Busses / Peripherals
- State specific observation focus
- Observation of mass-produced SoCs
- Unlimited time observation
- Low latency
- Intuitive to use

## Software Instrumentation


- + Easy to implement / low cost for tests
- Additional resources / different behavior
- Changing observation focus requires code recompilation
- Temporal assignment of different CPUs processes is very limited
- *Questionable approach: removing instrumentation from production code*

Original source code	Instrumented code (Statement / Condition Coverage)
<pre> void foo() {     bool found=false;     for (int i=0; (i&lt;100) &amp;&amp; (!found); ++i)     {         if (i==50) break;          if (i==20) found=true;          if (i==30) found=true;      }     printf("foo\n"); } </pre>	<pre> char inst[15]; void foo() {     bool found=false;     for (int i=0; ((i&lt;100)?inst[0]=1:inst[1]=1,0) &amp;&amp;         ((!found)?inst[2]=1:inst[3]=1,0); ++i)     {         if ((i==50?inst[4]=1:inst[5]=1,0))         { inst[6]=1; break;         }         if ((i==20?inst[7]=1:inst[8]=1,0))         { inst[9]=1; found=true;         }         if ((i==30?inst[10]=1:inst[11]=1,0))         { inst[12]=1; found=true;         }         inst[13]=1;     }     printf("foo\n");     inst[14]=1; } </pre>





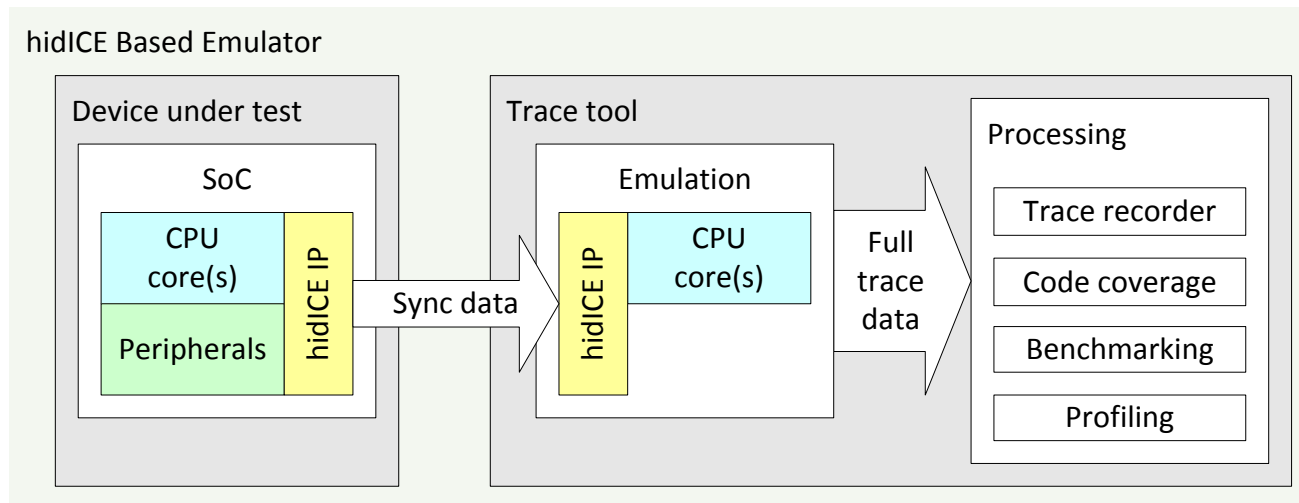
- Huang/Kao/Yang  
(National Sun Yat-Sen University Taiwan)
  - SYS-SIP SoC Development Infrastructure
  - three stages lossless instruction trace compression

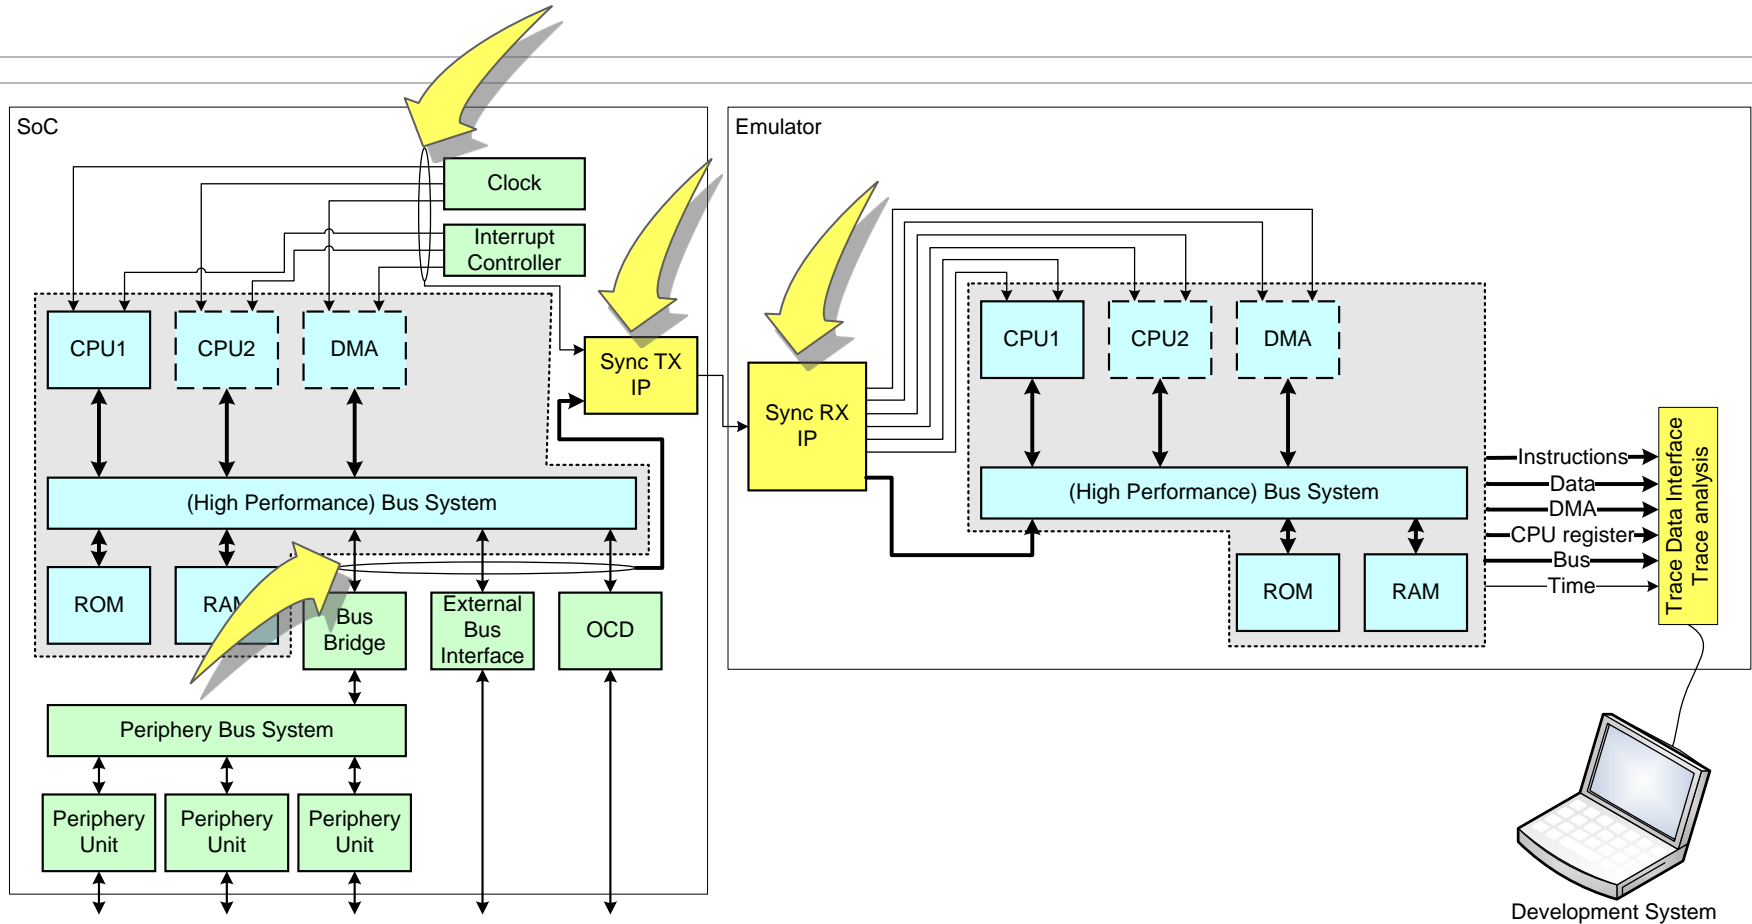


Fu-Ching Yang; , "SYS-SIP SoC Development Infrastructure", Dissertation, National Sun Yat-Sen University, 2009  
 Fu-Ching Yang; Yi-Ting Lin; Chung-Fu Kao; Ing-Jer Huang; , "An On-Chip AHB Bus Tracer With Real-Time Compression and Dynamic Multiresolution Supports for SoC", Very Large Scale Integration (VLSI) Systems, IEEE Transactions on , vol.19, no.4, pp.571-584, April 2011

## hidICE (*hidden ICE*)

- Emulate the SoC core region and access trace data from there
- Easy access to **full** trace data



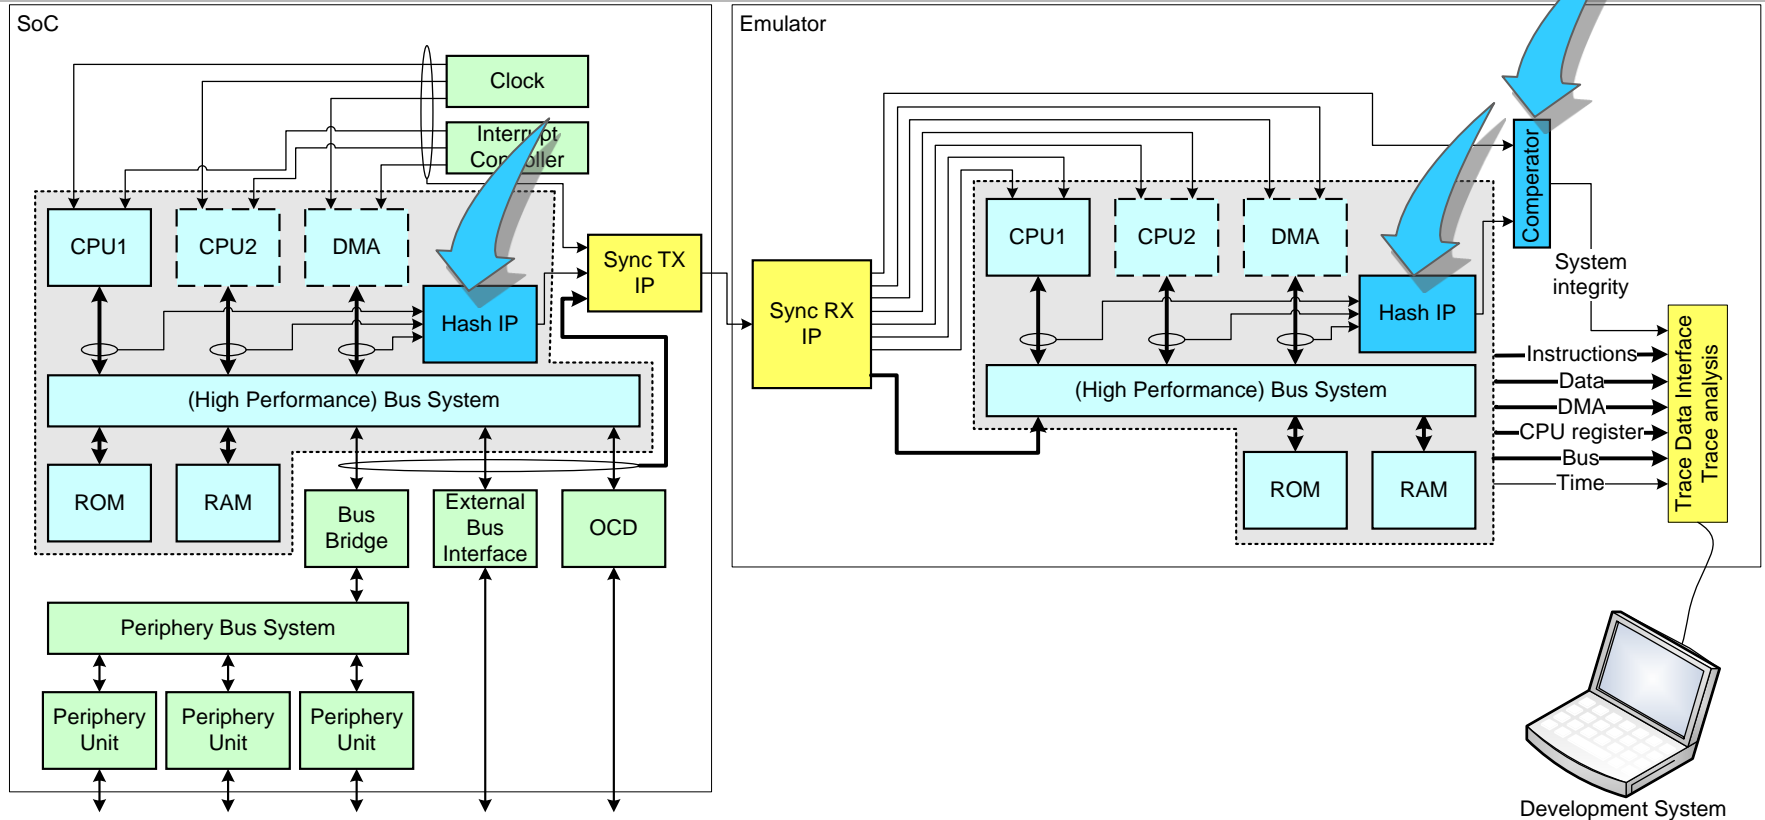


## Synchronization

Signals to transmit

Serialization

Deserialization



## Synchronization

Signals to transmit

Serialization

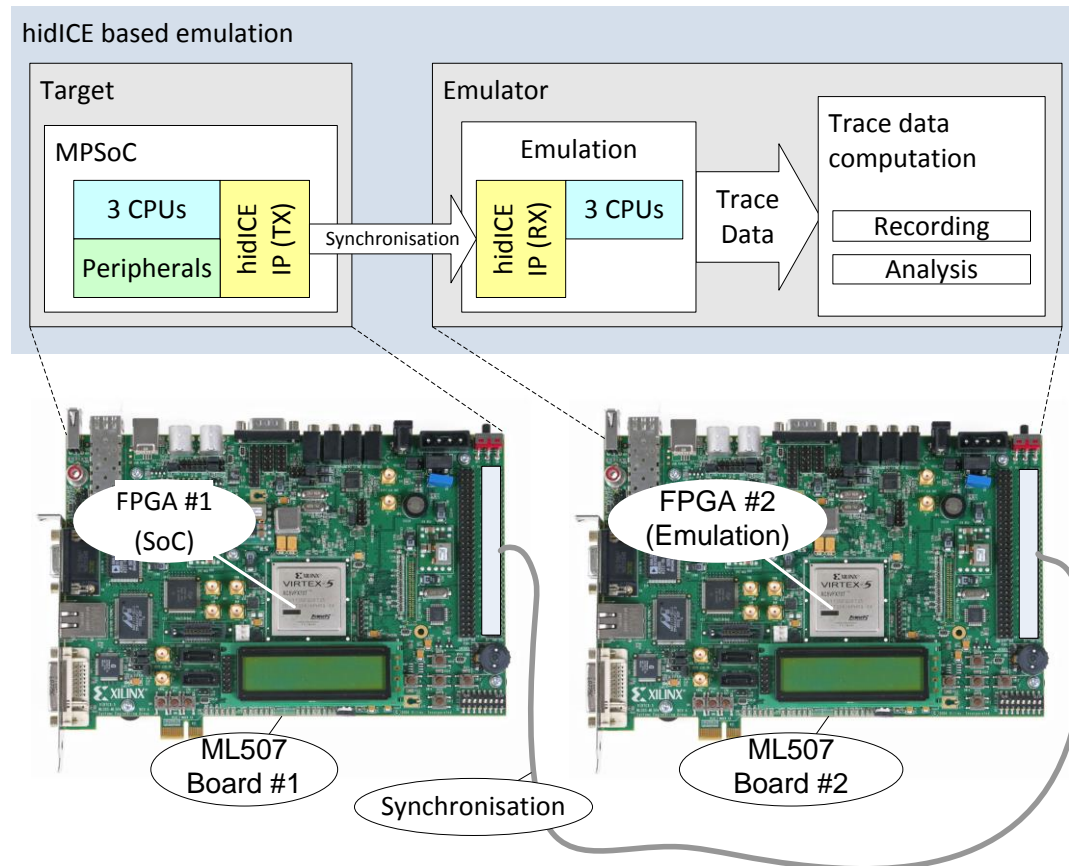
Deserialization

## System Integrity Control

Hash calculation

Hash check

## MPSoC implementation (3 x SPARC V8 / LEON3)



## Embedded trace

Time stamps	CPU 3
Instructions	
Data read	
Data written	
Time stamps	CPU 2
Instructions	
Data read	
Data written	
Time stamps	CPU 1
Instructions	
Data read	
Data written	
Time stamps	CPU 0
Instructions	
Data read	
Data written	

MPSoC (4 CPUs)

MPSoC, 4 CPUs, 1 GHz  
 Cycle accurate instruction + data trace  
 4 x 14 Bit x 1 GHz => approx. **28 Gbit/s**  
 ( + timestamps)  
 ( + bus master trace)  
 ( + peak bandwidth)

## hidICE

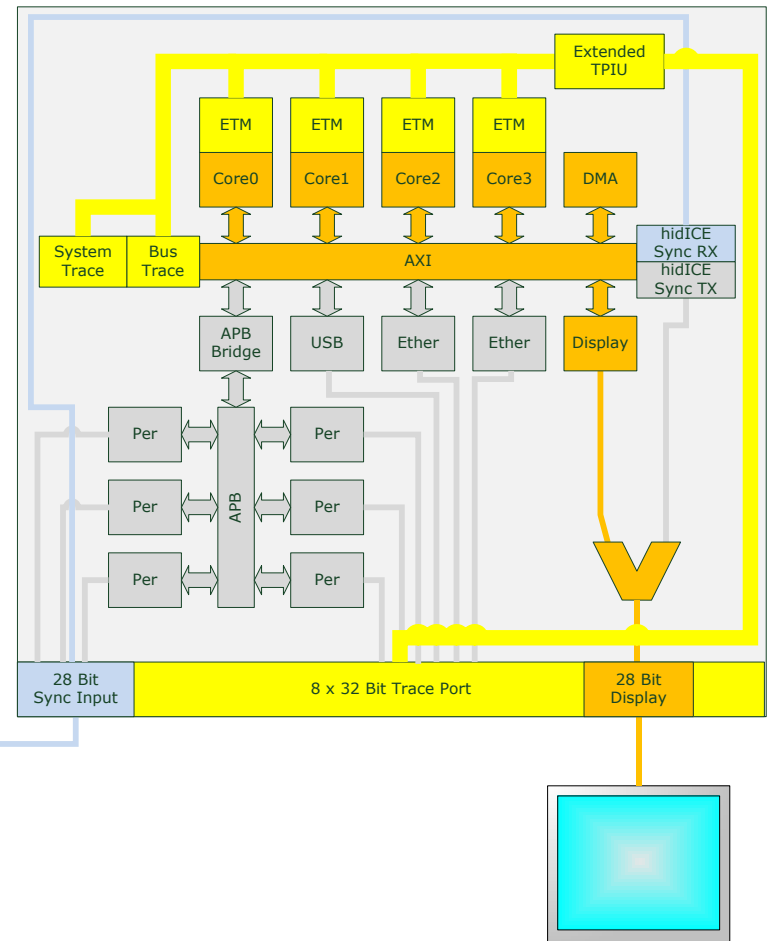
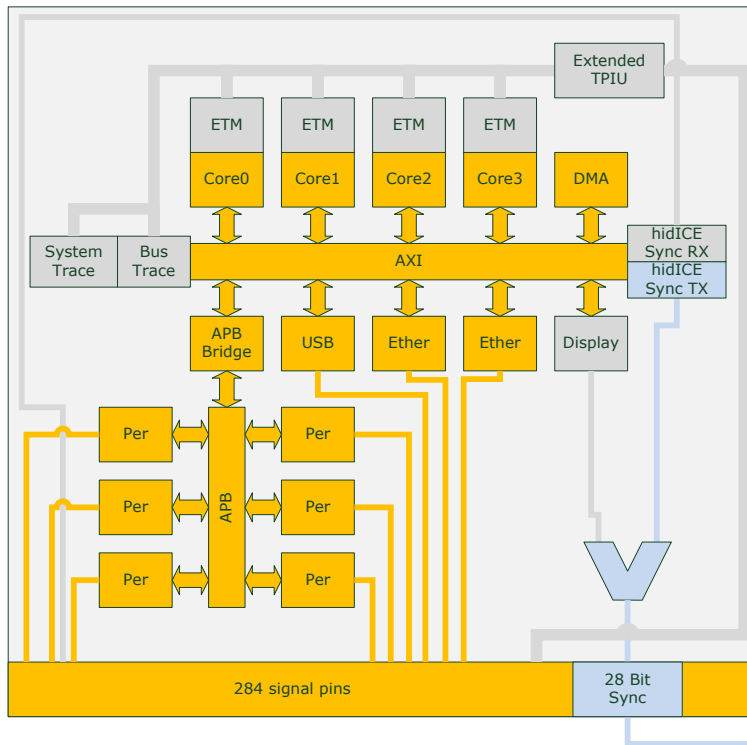
MPSoC, 4 CPUs, 1 GHz  
 1 x USB2.0, 2 x Gbit Ethernet,  
 some low speed peripherals

**Synchronization bandwidth: < 4 GBit**  
 - timestamps included  
 - bus master trace included  
 - peak bandwidth included

Events (CPU3)
Events (CPU2)
Events (CPU1)
Events (CPU0)
Data read (IO)

MPSoC (4 CPUs)

## Draft: hidICE for quad core SoC




## hidICE summary

- + Cycle accurate instruction and data trace from all CPUs
- + Cycle accurate data trace from all bus masters
- + Long-time observation
- + Low latency
- + Low intrusiveness (port replacement)
- Implementation effort (e.g. clock domain synchronization, correct implementation of the emulation)
- "All or none" – no partial trace
- Not applicable for SoCs with high I/O bandwidth

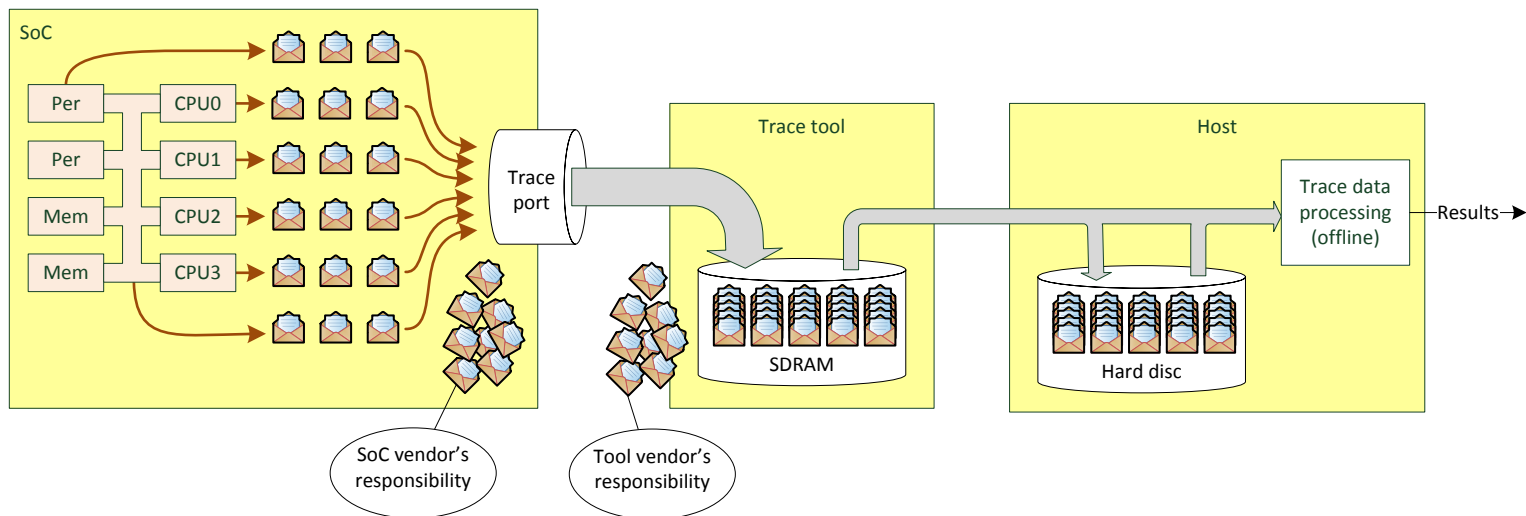


## Multi-Core observation challenges

- Make internal states visible
- Analyze trace data



## State of the Art: Trace data recording and offline processing


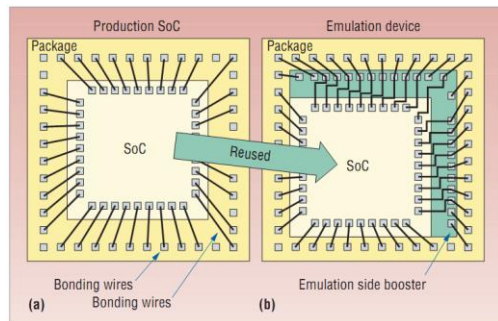


### Limitations:

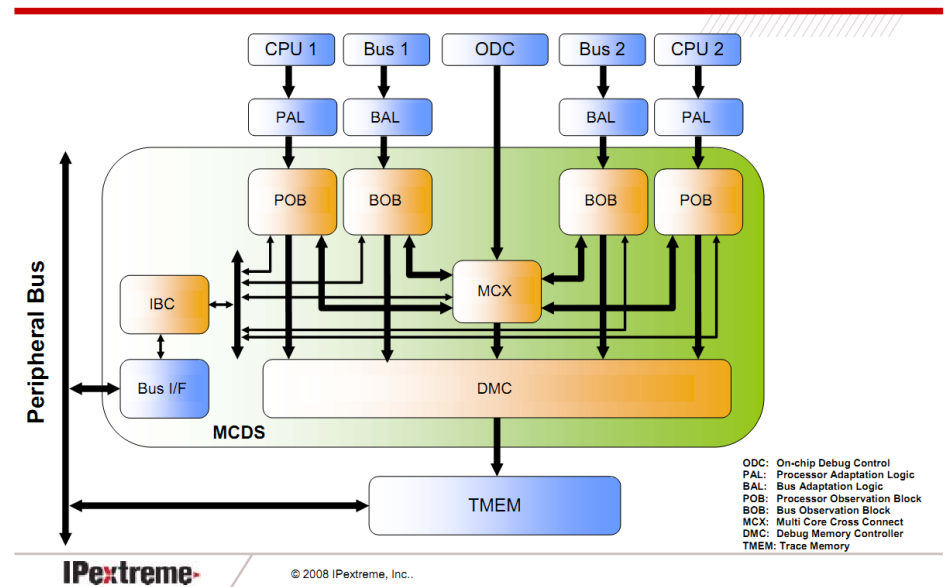
- Trace bandwidth  $\gg$  Processing bandwidth
- High latency in detection of SoC internal states state
- Limited in multiple observation focusses

## Infineon MCDS

- Combine SoC and trace tool



Mayer, A.; Siebert, H.; McDonald-Maier, K.D.; ,  
*"Boosting Debugging Support for Complex  
 Systems on Chip"* Computer , vol.40, no.4,  
 pp.76-81, April 2007

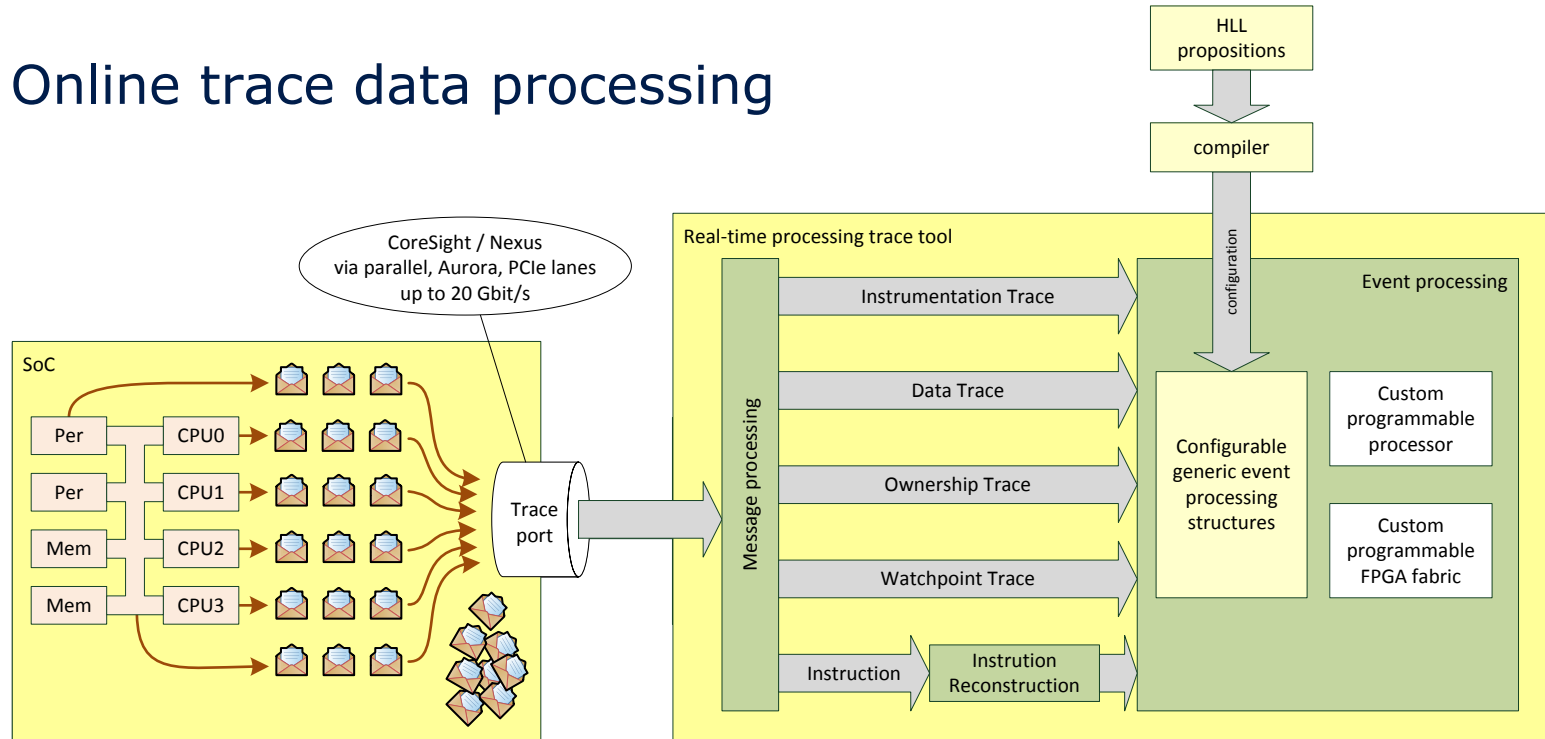


**IPextreme**

© 2008 IPextreme, Inc..

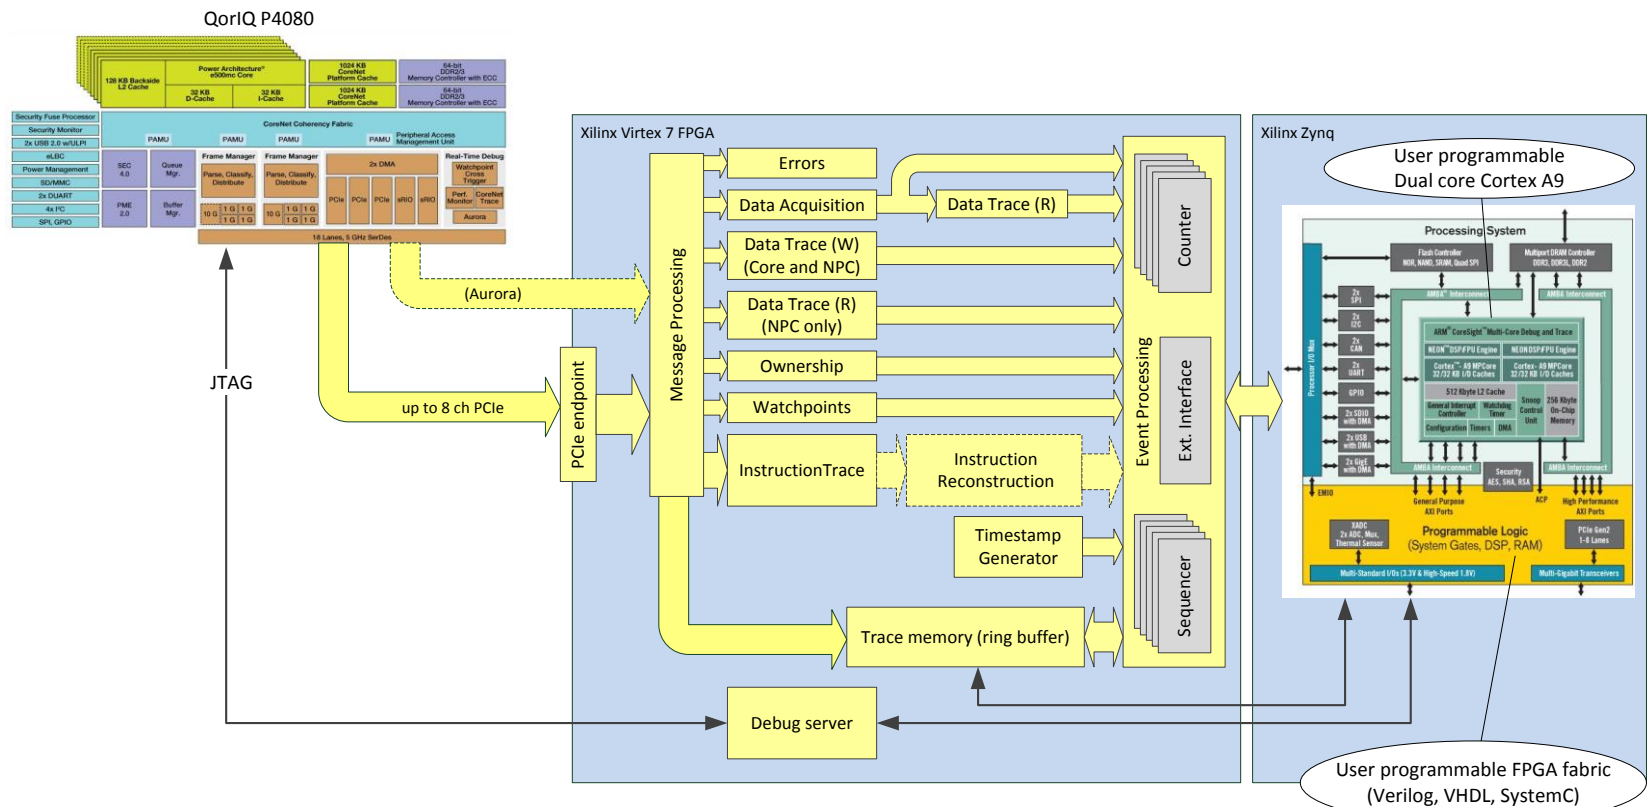
Source: [www.ipextreme.com](http://www.ipextreme.com)

## Online trace data processing



- Processing bandwidth  $\geq$  trace bandwidth
- $\mu$ s latency in detection of SoC internal states state
- Multiple observation focusses

## Online trace processing implementation example P4080 (8 cores e500mc @1.3 GHz)



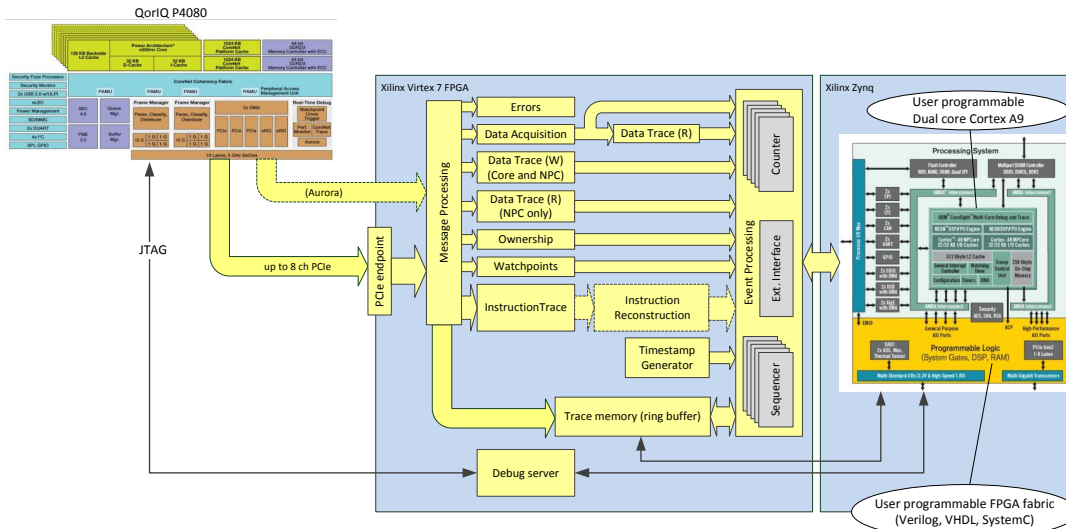
## Online trace processing

- High implementation effort (algorithms, FPGA resources) for
  - Transport protocol decoding (ARM CoreSight TPIU)
  - Detection of message boundaries (ARM CoreSight)
  - Online Reconstruction of direct branches (Nexus / ARM)
  - Runtime Verification infrastructure

Our wish to SOC vendors: More trace bandwidth!

## Online trace processing advantages

- + Real-time capability / unlimited observation time
- + Low latency of state specific changes of observation focus
- + Parallel observation of multiple hotspots
- + Low latency of state specific triggering of the SoC (e.g. tests)
- + Immediate access to the runtime information by the developer
- + Instruction / basic blocks level continuous profiling
- + Basic block level continuous WCET measurement
- + Automatic detection of race conditions (e.g. “Happened before” algorithm)
- + Extension of SoC internal debug / profiling structures
- + Combination of Trace debugging and Run/Stop debugging
- + The “swiss army knife” in debugging: Online Runtime Verification



Contact:  
 Alexander Weiss  
 Accemic GmbH & Co. KG  
 Franz-Huber-Str. 39  
 83088 Kiefersfelden

aweiss@accemic.com  
 +49 8033 6039790

In cooperation with



UNIVERSITÄT ZU LÜBECK